

Minimum spanning trees

A spanning tree of G is a subgraph T that is

- Connected
- Acyclic
- Includes all vertices

In: Connected undirected graph G with positive edge weights
 Out: Min weight spanning tree.

Greedy algorithm

Connected graph with distinct edge weights \Rightarrow MST exists and is unique

A cut in a graph is a partition of its vertices into two non-empty sets
 A crossing edge connects vertices in different sets.

Given any cut, the crossing edge of min weight is in the MST
 Proof in presentation

- Start with all edges colored gray
- Find cut with no black crossing edges; color its min-weight edge black.
- Repeat until $V-1$ edges are colored black. ($E = V-1$)

This computes the MST. Proof in presentation.

If the weights are not distinct, this algorithm still works.

If all edges are not connected, a single MST does not exist.

Weighted edge API

- Edge(int v, int w, double weight)
- int either()
- int other(int v)
- compareTo(Edge that)
- double weight()

Edge-weighted graph: vertex-indexed array of edges

Minimum spanning tree API

- MST(EdgeWeightedGraph G)
- Iterable<Edge> edges()
- double weight()

Kruskal's algorithm

- Consider edges in ascending order of weight.
- Add next edge to tree T unless doing so would create a cycle.
 Union-find can be used to detect cycles.
 (and should)

computes MST in time proportional to $E \log E$ (worst case)

Prim's algorithm

- Start with vertex 0 and greedily grow tree T
- Add to T the min weight edge with exactly one endpoint in T
- Repeat until $V-1$ edges

Lazy and eager implementation in presentation

Shortest path

- Single source, single sink, source-sink, all pairs.

Restrictions on edge weights

- Nonnegative weights
- Euclidean weights
- Arbitrary weights

Api

DirectedEdge(int v, int w, double weight)

int from()

int to()

double weight()

String toString()

Shortest-paths tree (SPT), represented by two arrays

- $\text{distTo}[v]$ is length of shortest path from s to v .
- $\text{edgeTo}[v]$ is last edge on shortest path from s to v

Relax edge $e = v \rightarrow w$

- $\text{distTo}[v]$ is length of shortest known path from s to v .
- $\text{distTo}[w]$ is length of shortest known path from s to w .
- $\text{edgeTo}[w]$ is last edge on shortest known path from s to w .
- If $e = v \rightarrow w$ gives shorter path to w through v ,
update both $\text{distTo}[w]$ and $\text{edgeTo}[w]$.

Let G be an edge-weighted digraph.

Then $\text{distTo}[]$ are the shortest path distances from s iff

- $\text{distTo}[s] = 0$
- For each vertex v , $\text{distTo}[v]$ is the length of some path from s to v .
- For each edge $e = v \rightarrow w$, $\text{distTo}[w] \leq \text{distTo}[v] + e.\text{weight}()$.

Proof in presentation

Generic algorithm (to compute SPT from s)

- Initialize $\text{distTo}[s] = 0$ and $\text{distTo}[v] = \infty$ for all other vertices.
- Repeat until optimality conditions are satisfied:
 - Relax any edge.

Dijkstra's algorithm

- Consider vertices in increasing order of distance from s .
(non-tree vertex with the lowest $\text{distTo}[]$ value)
- Add vertex to tree and relax all edges pointing from that vertex.

Acyclic shortest paths

- Consider vertices in topological order.
- Relax all edges pointing from that vertex.